

Lecture 6: Eligibility Traces

Lei Zhang

Outline

- Prediction: $TD(\lambda)$
- Control: $Sarsa(\lambda)$ & $Q(\lambda)$
- The Unified View of RL Solutions


*Materials are modified from David Silver's RL lecture notes

Outline

- Prediction: $TD(\lambda)$
- Control: Sarsa(λ) & $Q(\lambda)$
- The Unified View of RL Solutions

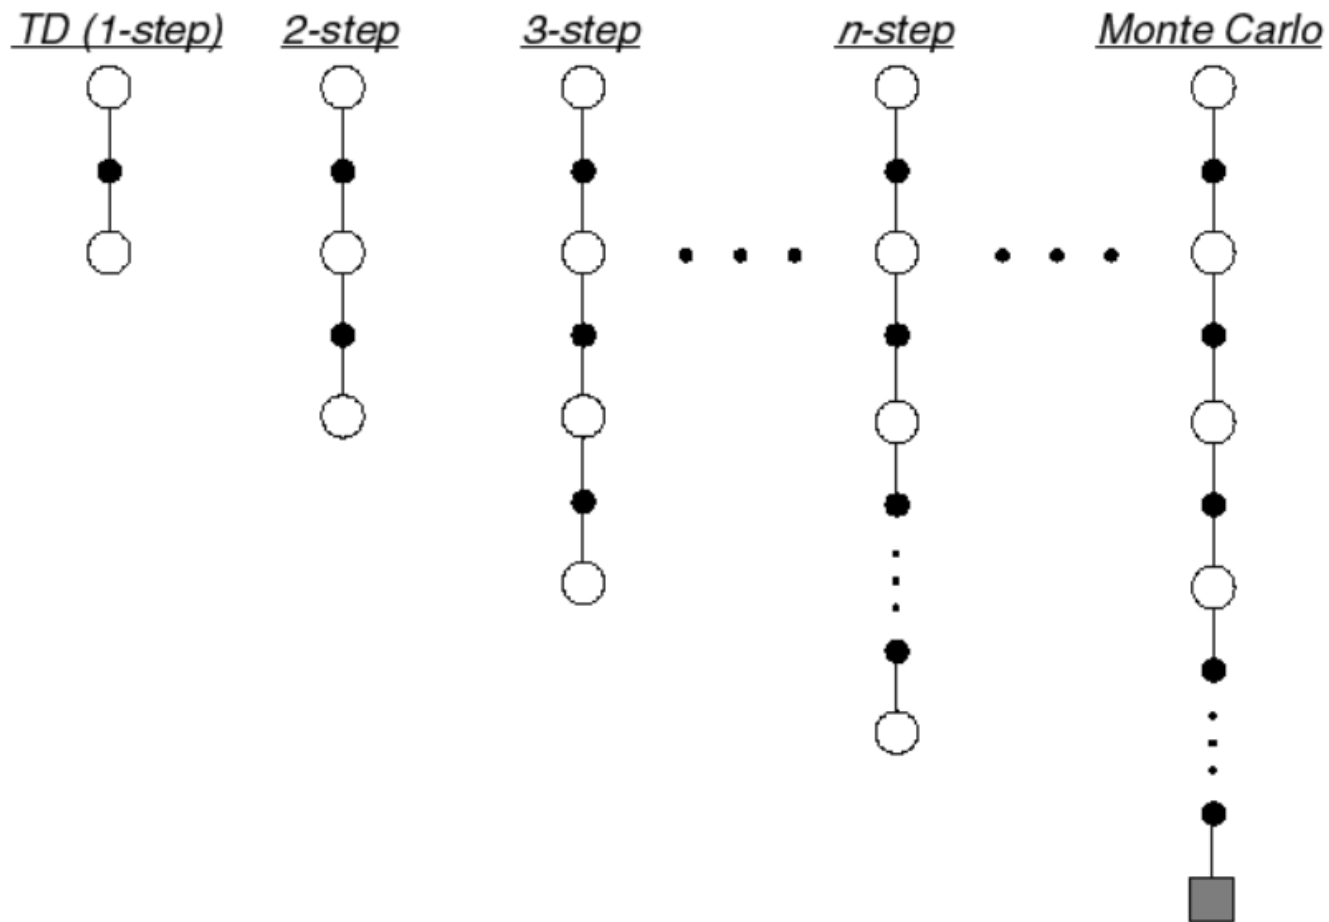
TD(0) Backup (Refresher)

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



n-step prediction

Let TD target look n steps into the future



n-step return

- Consider the following n -step returns for $n = 1, 2, \infty$:

$$\begin{array}{ll} n = 1 & (TD) \quad G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1}) \\ n = 2 & \quad \quad G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2}) \\ & \quad \quad \vdots \\ n = \infty & (MC) \quad G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T \end{array}$$

- Define the n -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

- n -step temporal-difference learning


$$V(S_t) \leftarrow V(S_t) + \alpha \left(G_t^{(n)} - V(S_t) \right)$$

Averaged n-step return

- We can average n -step returns over different n
e.g. average the 2-step and 4-step returns

$$\frac{1}{2}G^{(2)} + \frac{1}{2}G^{(4)}$$

- Combines information from two different time-steps
- Can we efficiently combine information from all time-steps?




λ -return

- The λ -return G_t^λ combines all n -step returns $G_t^{(n)}$
- Using weight $(1 - \lambda)\lambda^{n-1}$


$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

- Forward-view TD(λ)

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$




Weighting Function



$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

Forward View TD(λ)



- Update value function towards the λ -return
- Forward-view looks into the future to compute G_t^λ
- Like MC, can only be computed from complete episodes

Forward View & Backward View


- Forward view provides theory
- Backward view provides mechanism
- Update online, every step, from incomplete sequences

Eligibility Trace

- **Frequency heuristic**: assign credit to most frequent states
- **Recency heuristic**: assign credit to most recent states
- *Eligibility traces* combine both heuristics

$$E_0(s) = 0$$

$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$




Backward View TD(λ)

- Keep an eligibility trace for every state s
- Update value $V(s)$ for every state s
- In proportion to TD-error δ_t and eligibility trace $E_t(s)$

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



TD(λ) and TD(0)

- When $\lambda = 0$, only current state is updated

$$E_t(s) = \mathbf{1}(S_t = s)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

- This is exactly equivalent to TD(0) update

$$V(S_t) \leftarrow V(S_t) + \alpha \delta_t$$

Offline Equivalence of Forward and Backward Views

- Updates are accumulated within episode
- but applied in batch at the end of episode

The sum of offline updates is identical for forward-view and backward-view TD(λ)

$$\sum_{t=1}^T \alpha \delta_t E_t(s) = \sum_{t=1}^T \alpha \left(G_t^\lambda - V(S_t) \right) \mathbf{1}(S_t = s)$$

On-line TD(λ)

```
Initialize  $V(s)$  arbitrarily
Repeat (for each episode):
  Initialize  $Z(s) = 0$ , for all  $s \in \mathcal{S}$ 
  Initialize  $S$ 
  Repeat (for each step of episode):
     $A \leftarrow$  action given by  $\pi$  for  $S$ 
    Take action  $A$ , observe reward,  $R$ , and next state,  $S'$ 
     $\delta \leftarrow R + \gamma V(S') - V(S)$ 
     $Z(S) \leftarrow Z(S) + 1$ 
    For all  $s \in \mathcal{S}$ :
       $V(s) \leftarrow V(s) + \alpha \delta Z(s)$ 
       $Z(s) \leftarrow \gamma \lambda Z(s)$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```

Notation: $Z(s) = E(s)$

- No equivalence b/w Backward and Forward.
- Step size is sufficiently small \rightarrow almost “equivalence”
- Online TD updates are generally better than off-line TD.

Online Equivalence of Forward and Backward Views

- *True on-line TD(λ)* :
 - Van Seijen & Sutton, “True Online TD(λ)” ICML 2014
 - Van Seijen, etc “True Online TD Learning”, JMLR, 2016

- Forward View: truncated λ -return

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{t'-1} \lambda^{n-1} G_t^{(n)} + \lambda^{t'-1} G_t^{(t')}$$

- Backward View: a slightly different form of eligibility trace

Summary

| Offline updates | $\lambda = 0$ | $\lambda \in (0, 1)$ | $\lambda = 1$ |
|-----------------|---------------|---------------------------------|-----------------------|
| Backward view | TD(0) | TD(λ) | TD(1) |
| Forward view | TD(0) | Forward TD(λ) | MC |
| Online updates | $\lambda = 0$ | $\lambda \in (0, 1)$ | $\lambda = 1$ |
| Backward view | TD(0) | TD(λ) \nparallel | TD(1) \nparallel |
| Forward view | TD(0) | Forward TD(λ) | MC |
| Exact Online | TD(0) | Exact Online TD(λ) | Exact Online TD(1) |

Proof?

= here indicates equivalence in total update at end of episode.

Outline

- Prediction: $TD(\lambda)$
- Control: $Sarsa(\lambda)$ & $Q(\lambda)$
- The Unified View of RL Solutions

n-step Sarsa

- Consider the following n -step returns for $n = 1, 2, \infty$:

$$\begin{array}{ll} n = 1 & \text{(Sarsa)} \quad q_t^{(1)} = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) \\ n = 2 & q_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 Q(S_{t+2}, A_{t+2}) \\ & \vdots \\ n = \infty & \text{(MC)} \quad q_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T \end{array}$$

- Define the n -step Q-return

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n}, A_{t+n})$$

- n -step Sarsa updates $Q(s, a)$ towards the n -step Q-return

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left(q_t^{(n)} - Q(S_t, A_t) \right)$$


Forward View Sarsa(λ)

- The q^λ return combines all n -step Q-returns $q_t^{(n)}$
- Using weight $(1 - \lambda)\lambda^{n-1}$

$$q_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)}$$

- Forward-view Sarsa(λ)

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left(q_t^\lambda - Q(S_t, A_t) \right) \quad \sum = 1$$



Backward View Sarsa(λ)

- Just like TD(λ), we use **eligibility traces** in an online algorithm
- But Sarsa(λ) has one eligibility trace for each state-action pair

$$E_0(s, a) = 0$$

$$E_t(s, a) = \gamma\lambda E_{t-1}(s, a) + \mathbf{1}(S_t = s, A_t = a)$$

- $Q(s, a)$ is updated for every state s and action a
- In proportion to TD-error δ_t and eligibility trace $E_t(s, a)$

$$\delta_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \delta_t E_t(s, a)$$


Sarsa (λ) Algorithm

```
Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
Repeat (for each episode):
     $Z(s, a) = 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
    Initialize  $S, A$ 
    Repeat (for each step of episode):
        Take action  $A$ , observe  $R, S'$ 
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
         $\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$ 
         $Z(S, A) \leftarrow Z(S, A) + 1$ 
        For all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ :
             $Q(s, a) \leftarrow Q(s, a) + \alpha \delta Z(s, a)$ 
             $Z(s, a) \leftarrow \gamma \lambda Z(s, a)$ 
         $S \leftarrow S'; A \leftarrow A'$ 
    until  $S$  is terminal
```


Notation: $Z(s, a) = E(s, a)$

Sarsa (λ) Gridworld Example


Path taken




Action values increased by one-step Sarsa



Action values increased
by 10-step Sarsa



Action values increased by Sarsa(λ) with $\lambda=0.9$

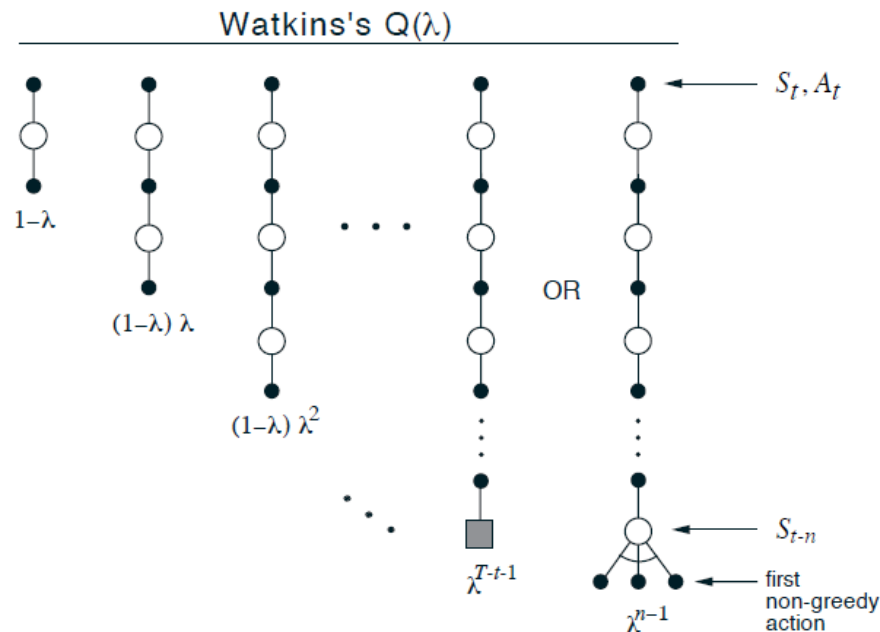


Q (λ) Algorithm – Forward View

- We show Watkins's Q(λ). See literature for others Q(λ) algorithms
- n-step Q return: lookahead stops at the first non-greedy action (i.e. exploratory action). If action at $t+n$ is the first non-greedy action, then the longest backup is toward

$$R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n \max_a Q_t(S_{t+n}, a)$$

- Backup diagram



Q (λ) Algorithm – Backward View

- One eligibility trace for each state-action pair

$$E_t(s, a) = \mathbf{1}(S_t = s, A_t = a) + \gamma\lambda E_{t-1}(s, a) \times \mathbf{1}(Q_{t-1}(S_t, A_t) = \max_a Q_{t-1}(S_t, a))$$

- Q update

$$\delta_t = R_{t+1} + \gamma \max_{a'} Q_t(S_{t+1}, a') - Q_t(S_t, A_t)$$

$$Q_{t+1}(s, a) = Q_t(s, a) + \alpha \delta_t E_t(s, a)$$

Q (λ) Algorithm

```
Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
Repeat (for each episode):
   $Z(s, a) = 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
  Initialize  $S, A$ 
  Repeat (for each step of episode):
    Take action  $A$ , observe  $R, S'$ 
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $A^* \leftarrow \arg \max_a Q(S', a)$  (if  $A'$  ties for the max, then  $A^* \leftarrow A'$ )
     $\delta \leftarrow R + \gamma Q(S', A^*) - Q(S, A)$ 
     $Z(S, A) \leftarrow Z(S, A) + 1$ 
    For all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ :
       $Q(s, a) \leftarrow Q(s, a) + \alpha \delta Z(s, a)$ 
      If  $A' = A^*$ , then  $Z(s, a) \leftarrow \gamma \lambda Z(s, a)$ 
      else  $Z(s, a) \leftarrow 0$ 
     $S \leftarrow S'; A \leftarrow A'$ 
  until  $S$  is terminal
```

Notation: $Z(s, a) = E(s, a)$

Extensions

- Disadvantage of Watkins's $Q(\lambda)$?
- Other $Q(\lambda)$ algorithms
 - Peng's $Q(\lambda)$
 - Naïve $Q(\lambda)$

Outline

- Prediction: $TD(\lambda)$
- Control: $Sarsa(\lambda)$ & $Q(\lambda)$
- The Unified View of RL Solutions

The Unified View of RL Solutions

